

Tools and Techniques, Lecture 3

Performance

- Algorithmic complexity
- Memory access patterns
- Benchmarking



LUNDS UNIVERSITET



CERN
School of Computing

CERN School of Computing 2025, Lund
Sten Åstrand, Lund University

Why?



Image credit: Giulio Eulisse, "Benchmarking and Profiling" lecture, CSC2024

How does it scale?

ALICE
Event Processing
Nodes Cluster

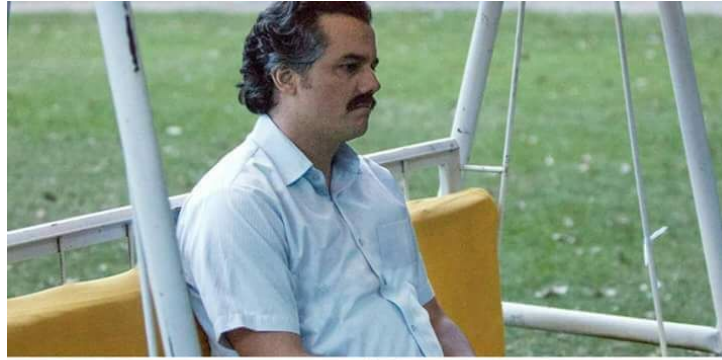
24640 CPU cores
2800 GPUs



Image credit: Giulio Eulisse, "Benchmarking and Profiling" lecture, CSC2024



How do you scale?



*Man making Python plots,
color photograph*



Image from imgflip.com

Algorithmic complexity

Theoretical metric for estimating resource consumption (in our case runtime - time complexity).

“The time complexity of an algorithm represents the number of steps it has to take to complete.”

– Complexity Theory. Brilliant.org. Retrieved 13:58, June 18, 2025,
from <https://brilliant.org/wiki/complexity-theory/>

Also useful: Devopedia. 2022. "Algorithmic Complexity." Version 8, February 19. Accessed 2024-06-25.
<https://devopedia.org/algorithmic-complexity>

Big O notation




⇒ algorithm is on the order of $f(n)$ or $O(f(n))$

Example: if an algorithm's runtime grows **linearly with the input size**, the **algorithm is $O(n)$**


Strictly: $O(f(n)) \Leftrightarrow$ “asymptotically bounded by $f(n)$ up to some constant”
(see Bachmann-Landau notation)

Big O Notation. Brilliant.org. Retrieved 13:54, June 18, 2025, from <https://brilliant.org/wiki/big-o-notation/>

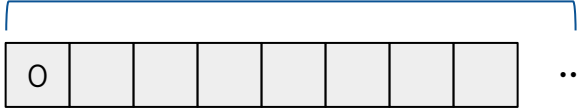
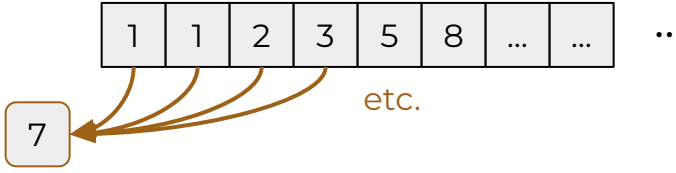
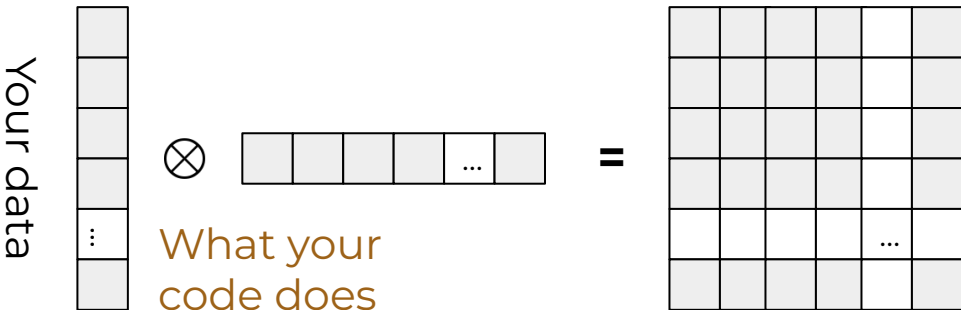
Common Big O's and examples thereof




Common Big O's and examples thereof






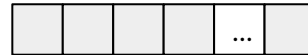
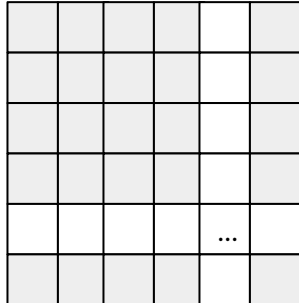
Common Big O's and examples thereof

	n elements	Complexity	Time
Your data		$O(1)$	Constant
Your data What your code does		?	?
Your data What your code does		?	?


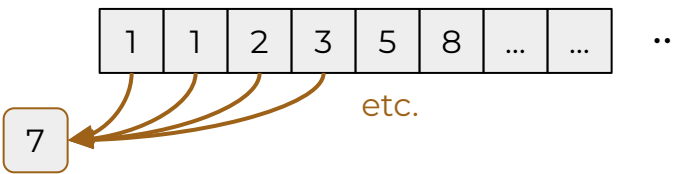
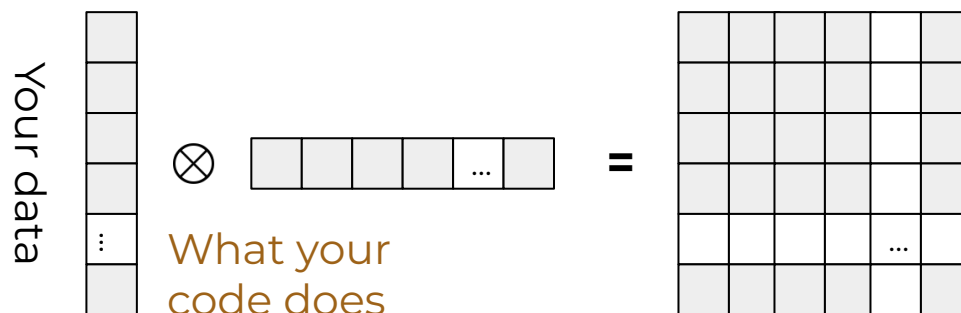
Common Big O's and examples thereof

	n elements	Complexity	Time
Your data		$O(1)$	Constant

Your data			
What your code does		$O(n)$	Linear

Your data		\otimes 	=		?	?
		What your code does				

Common Big O's and examples thereof

	n elements	Complexity	Time
Your data		$O(1)$	Constant
Your data What your code does		$O(n)$	Linear
Your data What your code does		$O(n^2)$	Quadratic

Common Big O's and examples thereof

$O(n)$ - linear time (i.e. input size doubles = runtime doubles)

- Radix sort

$O(n^2)$ - quadratic time (input size doubles = runtimes quadruples)

- “For loop in a for loop”
- Outer product of two n -length vectors $\Rightarrow n$ -by- n matrix
- Selection sort, insertion sort

Common Big O's and examples thereof

$O(n)$ - linear time (i.e. input size doubles = runtime doubles)

- Radix sort

aktshually $O(k \cdot n)$

$O(n^2)$ - quadratic time (input size doubles = runtimes quadruples)

- “For loop in a for loop”
- Outer product of two n -length vectors $\Rightarrow n$ -by- n matrix
- Selection sort, insertion sort

Common Big O's and examples thereof, cont.

$O(\log_2 n)$ - logarithmic time

- Binary search of a sorted array

$O(n \log_2 n)$ - log-linear or “linearithmic” time
(close to $O(n)$ for “reasonably sized” n)

- Many sorting algorithms, e.g. merge sort, heap sort

$O(n^3)$ - cubic time

- “For loop in a for loop in a for loop”
- “Schoolbook” matrix multiplication



Common Big O's and examples thereof, cont.

$O(\log_2 n)$ - logarithmic time

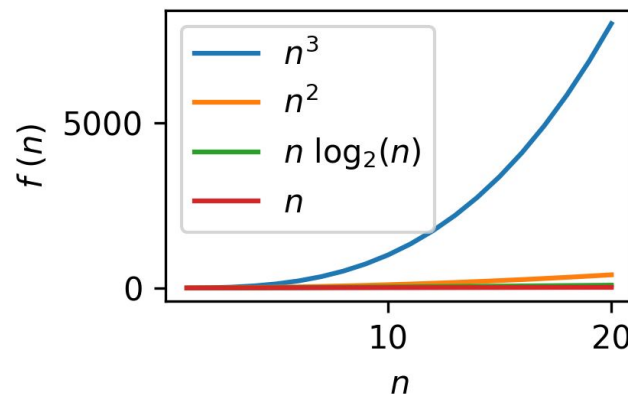
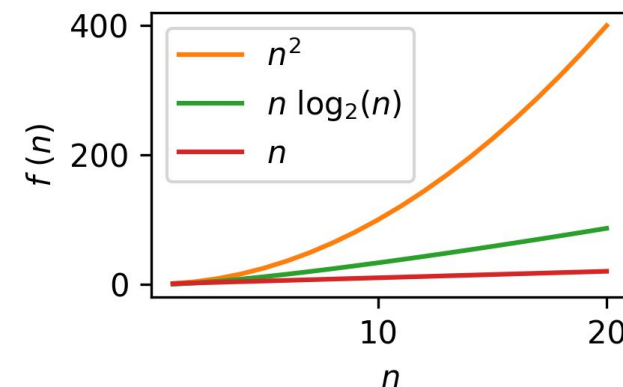
- Binary search of a sorted array

$O(n \log_2 n)$ - log-linear or “linearithmic” time
(close to $O(n)$ for “reasonably sized” n)

- Many sorting algorithms, e.g. merge sort, heap sort

$O(n^3)$ - cubic time

- “For loop in a for loop in a for loop”
- “Schoolbook” matrix multiplication



Two really Big O's (to avoid if possible)

$O(2^n)$ and **$O(n!)$** - exponential time and factorial time

- Brute-force search, combinatorics, NP-hard problems
- Finding prime numbers



Matrix multiplication

- “Schoolbook” matrix multiplication $O(n^3)$
- The Strassen algorithm (1969) $O(n^{2.805})$
- State of the art (2023) $O(n^{2.373})$

For 4 x 4 matrices:

Strassen algorithm: 49 multiplications

Google Deepmind AlphaEvolve algorithm: 48 multiplications

AlphaEvolve: A Gemini-powered coding agent for designing advanced algorithms

<https://deepmind.google/discover/blog/alphaevolve-a-gemini-powered-coding-agent-for-designing-advanced-algorithms/>





Algorithmic complexity - conclusion

```
int binary_search(const std::vector<int>& v, int target) {  
    int left = 0, right = v.size() - 1;  
  
    while (left <= right) {  
        int mid = left + (right - left) / 2;  
        if (v[mid] == target)  
            return mid;  
        else if (v[mid] < target)  
            left = mid + 1;  
        else  
            right = mid - 1;  
    }  
  
    return -1; // not found  
}
```

Common code patterns

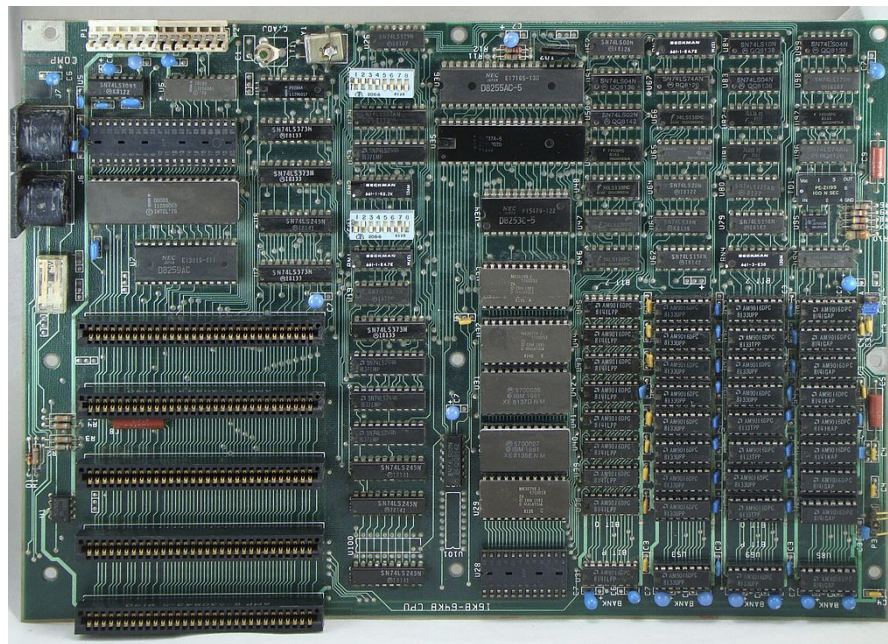
- For loop in a for loop
 - Solve differently?
 - Combine loops?
- Searching through data
 - Consider sorting first?
 - Use look-up table?
- Recomputing values
 - Compute once, keep in a table?



Pit stop



Memory access patterns



An IBM Personal Computer Model 5150 motherboard, 1981 - the first "PC". Image credit: user GermanX on Wikimedia Commons, under license [Attribution-ShareAlike 2.5 Generic](https://creativecommons.org/licenses/by-sa/2.5/)

Memory access patterns

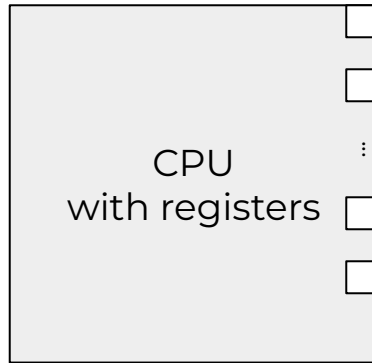
Image credit: user Gravislizard on
Wikimedia Commons, under license
[Attribution-ShareAlike 2.0 France](#)



An IBM Personal Computer Model 5150 motherboard, 1981 - the first "PC". Image credit: user GermanX on Wikimedia Commons, under license [Attribution-ShareAlike 2.5 Generic](#)

CPU-to-memory layout

CPU registers

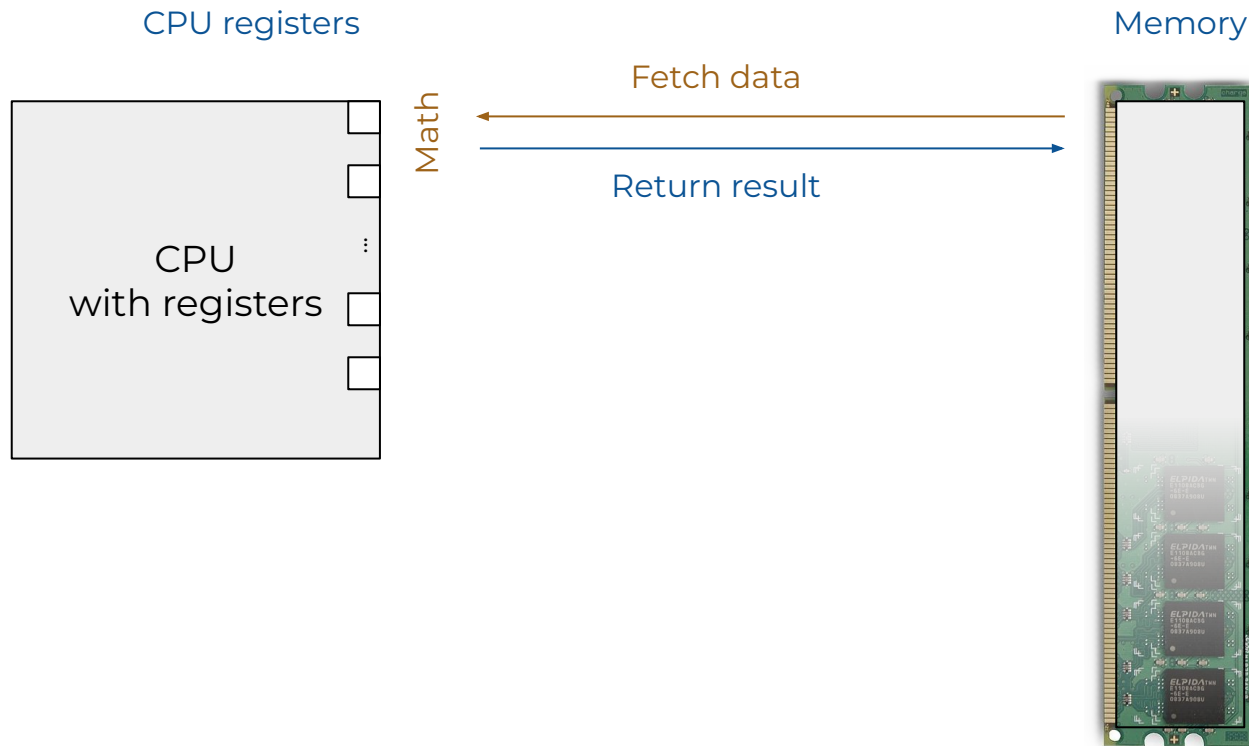


Memory



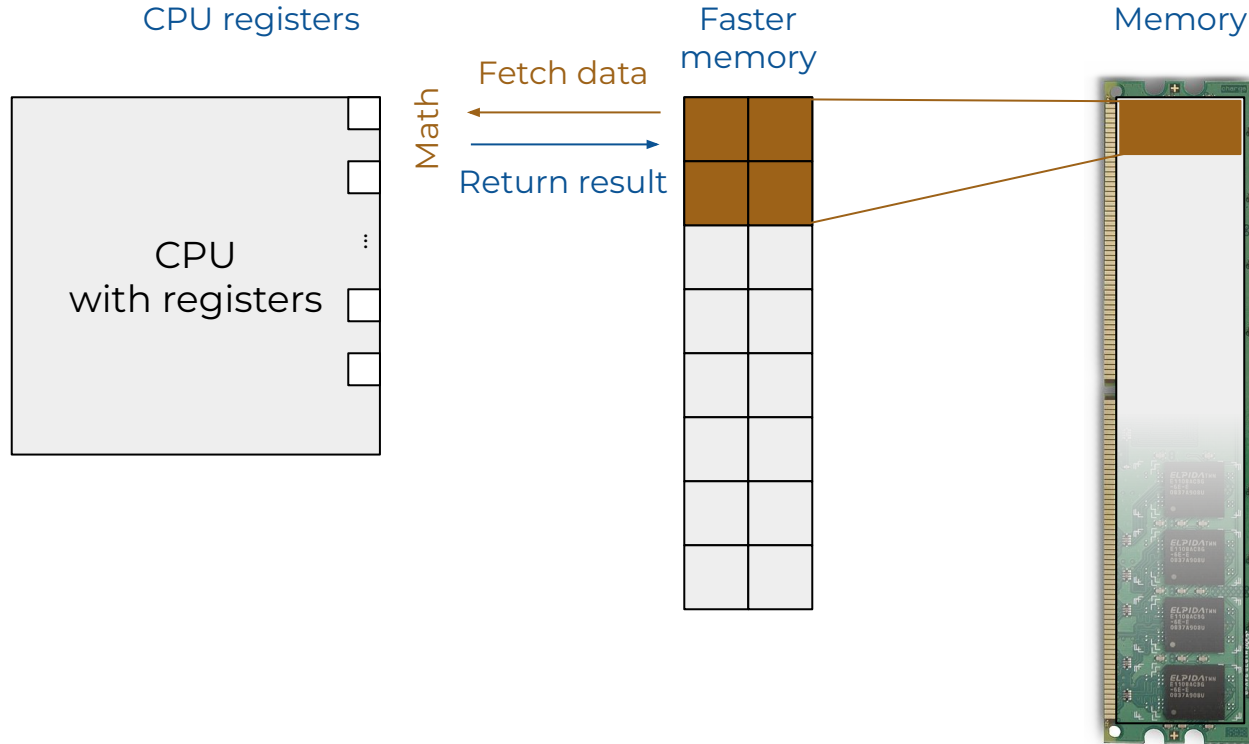
RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

CPU-to-memory layout



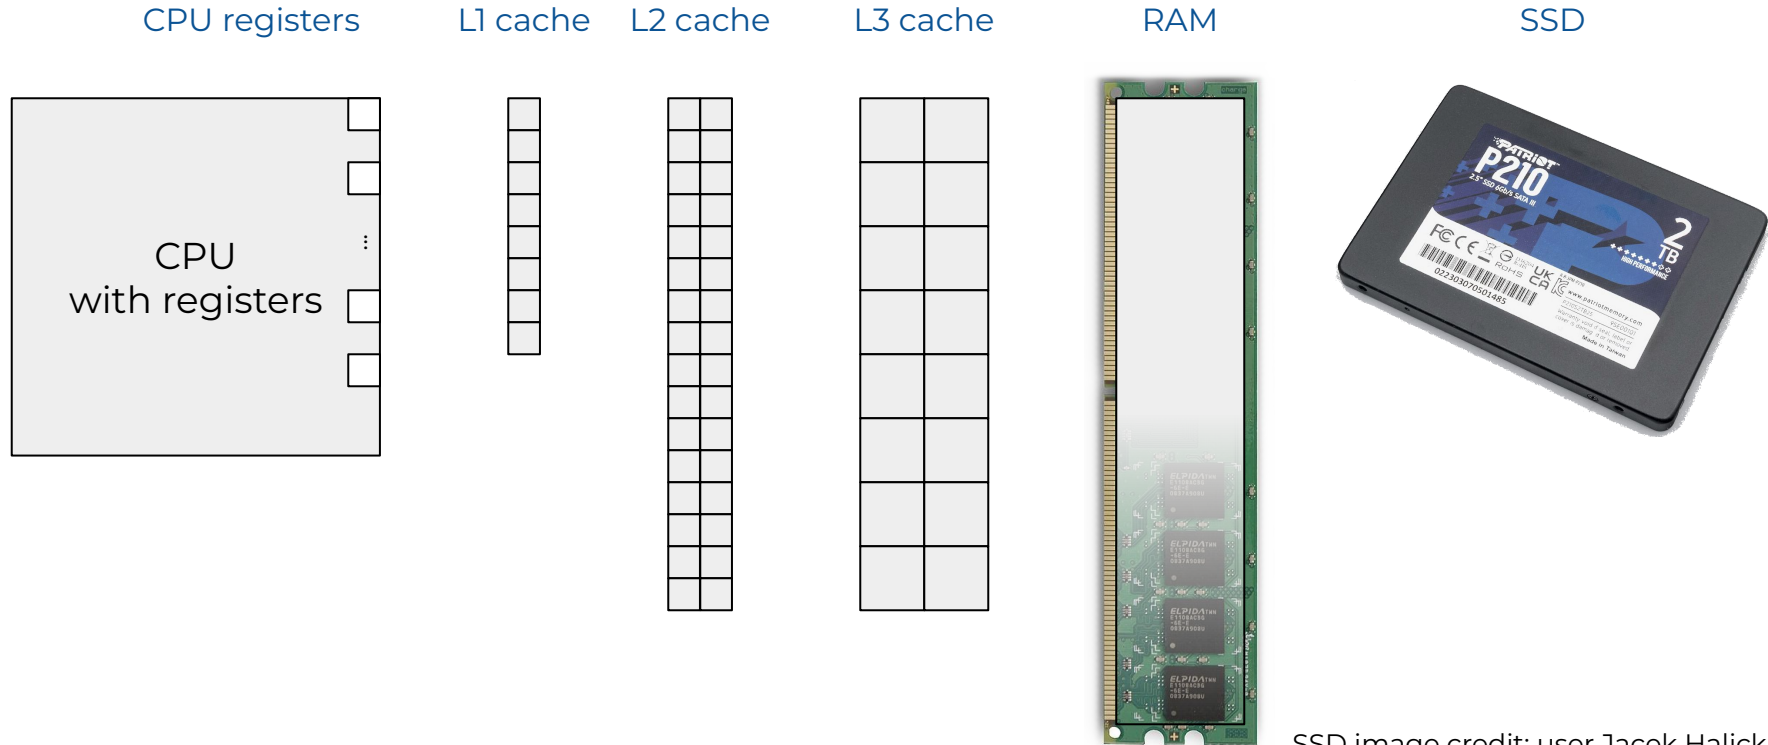
RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

CPU-to-memory layout



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

CPU-to-memory layout




RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

CPU-to-memory layout

Typical sizes of different memory hardware



RAM
O(10 GB)



SSD
O(1 TB)




RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

CPU-to-memory layout

Typical sizes of different memory hardware



RAM
O(10 GB)



100ns

SSD
O(1 TB)




16μs

RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

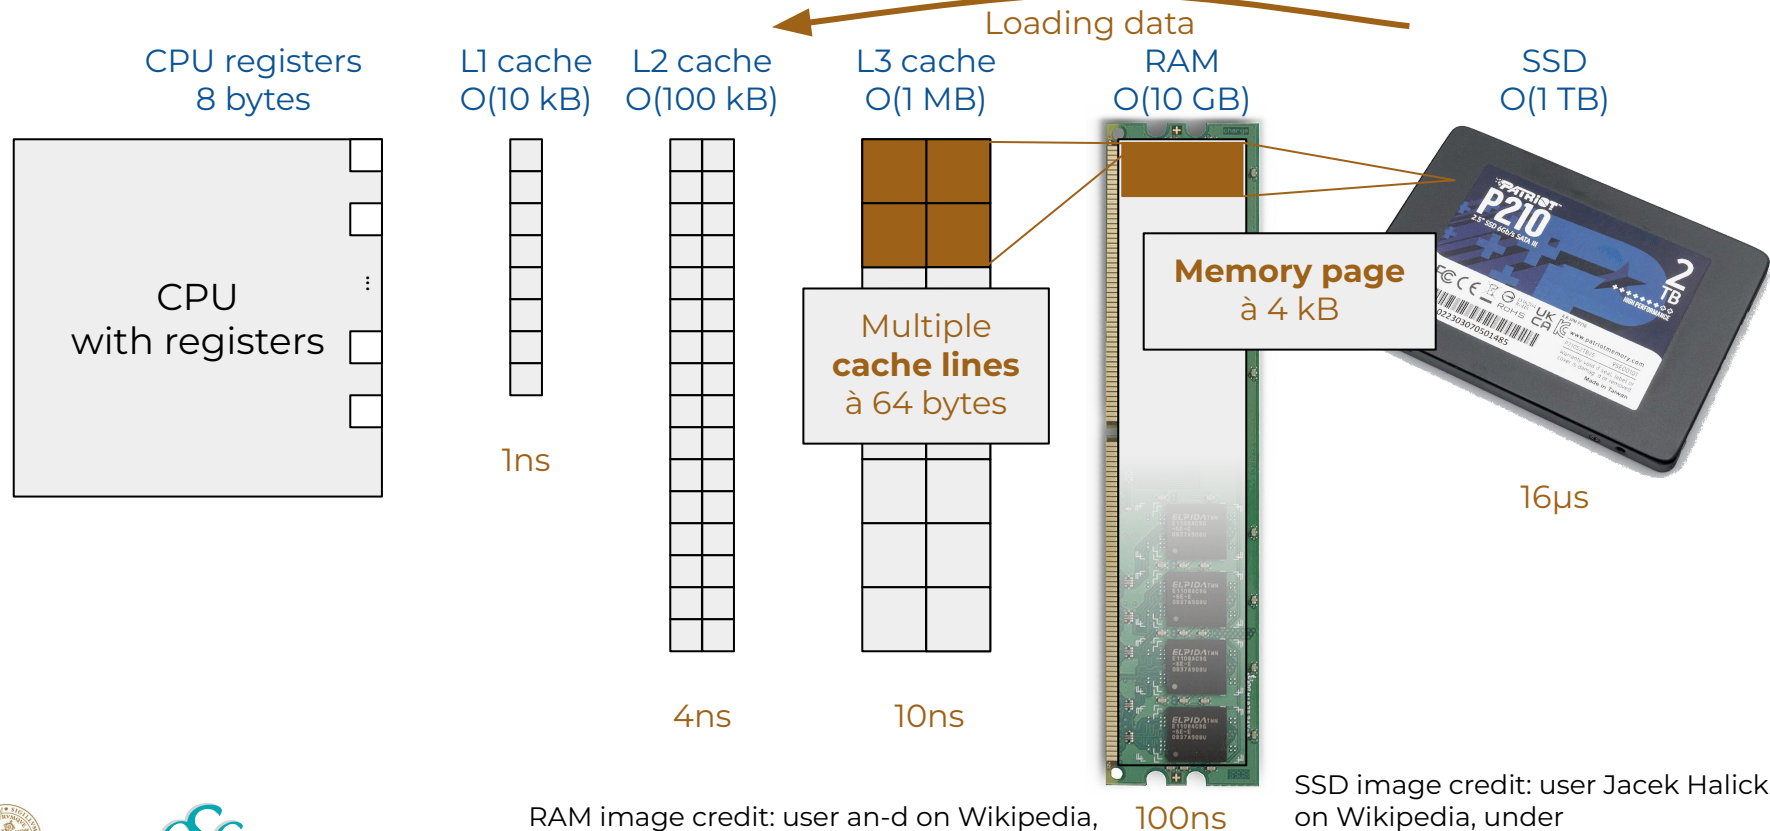
Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)


Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

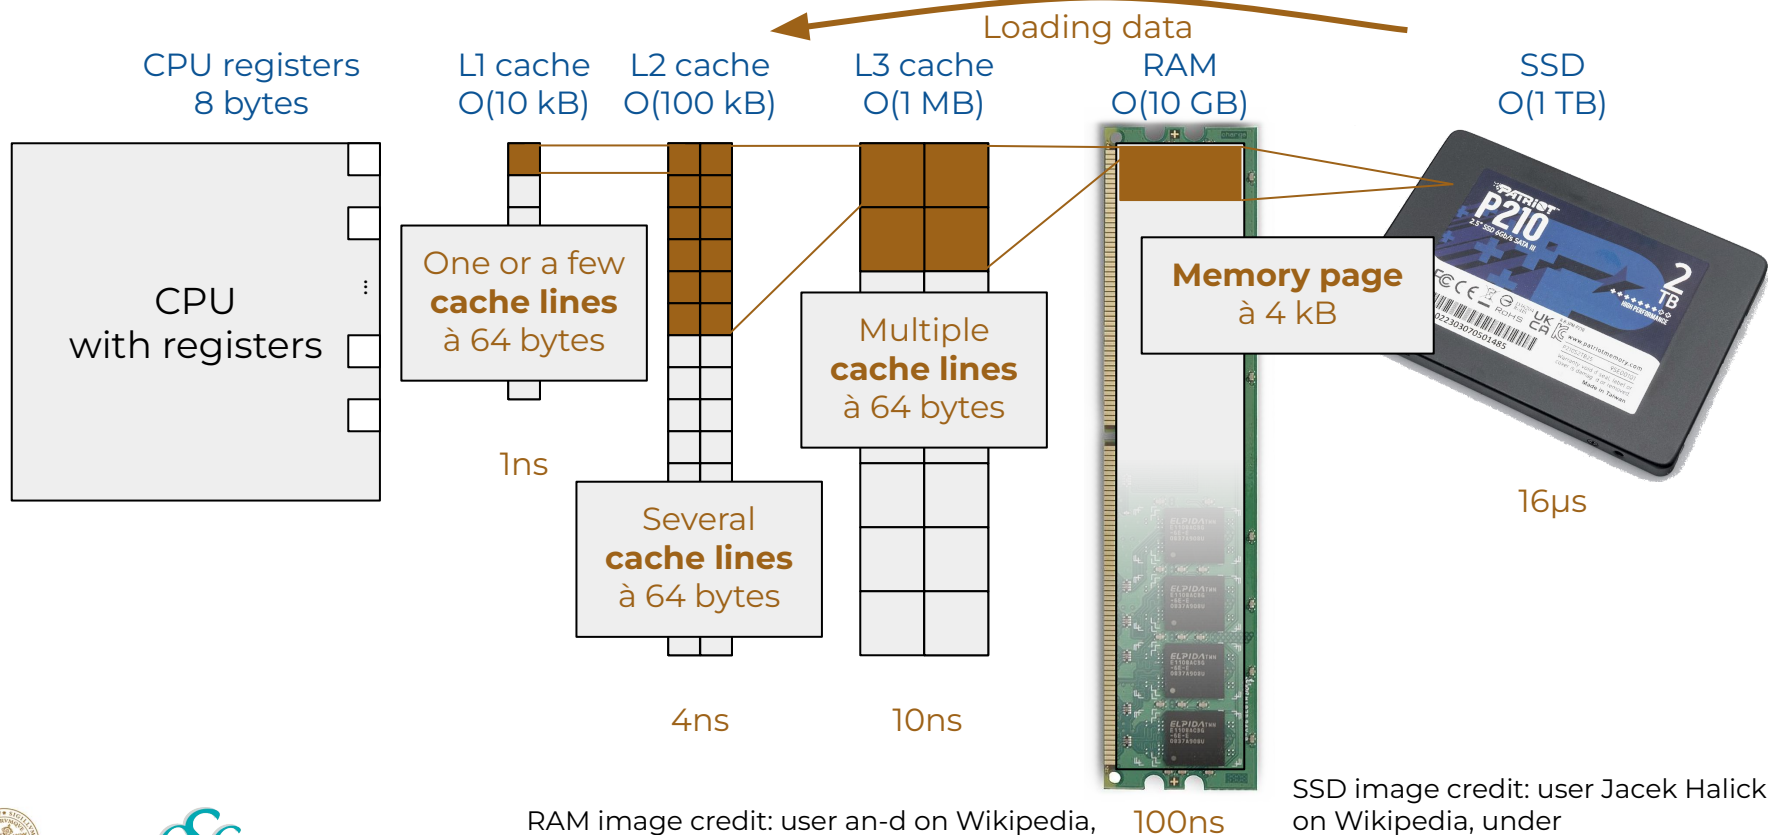
Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)


Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

Block-based memory access, the “why”

“**Principle of locality**” or “**data locality**” - data access often happens on many elements close to each other.

If you read
this...



You will
likely then
read those

Block-based memory access, the “why”

“**Principle of locality**” or “**data locality**” - data access often happens on many elements close to each other.

If you read
this...



You will
likely then
read those


- **space locality**
- **time locality**

Pre-fetching: loading data that is likely to be needed soon

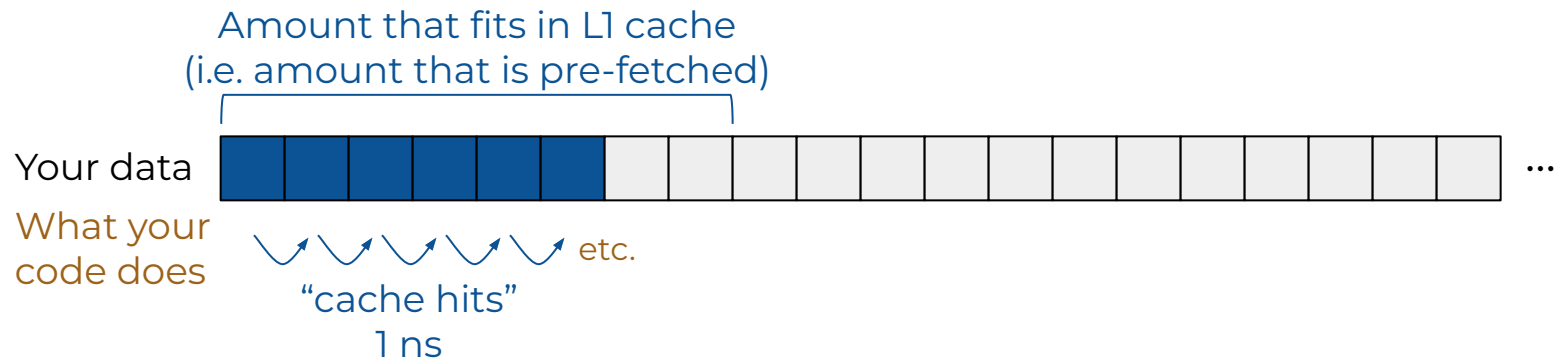
Memory access patterns

Amount that fits in L1 cache
(i.e. amount that is pre-fetched)


Your data



Memory access patterns



Memory access patterns



Memory access patterns

Amount that fits in L1 cache
(i.e. amount that is pre-fetched)

Your data

A horizontal bar representing memory. The first 8 segments are dark blue, followed by 1 brown segment, and then 10 light gray segments. An ellipsis follows. A bracket above the first 8 segments is labeled "Amount that fits in L1 cache (i.e. amount that is pre-fetched)".

What your
code does



"cache hits"

1 ns

"cache miss"

4 ns

Your data

A horizontal bar representing memory. The segments are: dark blue, light gray, light gray, light gray, light gray, dark blue, light gray, light gray, light gray, light gray, brown, light gray, light gray, light gray, light gray, dark blue, light gray, light gray, light gray, light gray. An ellipsis follows. A bracket above the first 10 segments is labeled "Amount that fits in L1 cache (i.e. amount that is pre-fetched)".

What your
code does



etc.

Memory access patterns

Amount that fits in L1 cache
(i.e. amount that is pre-fetched)

Your data



What your
code does




"cache hits"

1 ns


"cache miss"

4 ns

Your data




What your
code does



40




Memory access patterns - conclusions



Pit stop 2



Benchmarking



Benchmarking, pitfalls of

...but measuring is hard.

- Compiler optimizations
- Programs running in parallel \Rightarrow variance
- CPUs boost clock frequencies
- ...

Takeaway: healthy scepticism and reasonable expectations



Linux **time** command

```
int main() {  
    int temp;  
    for (int i = 0; i < 1000000000; i++) {  
        temp = i;  
    }  
    return temp;  
}
```

```
$ g++ -O3 time_example.cpp -o time_example  
$ time ./time_example
```

```
real    0m0.009s  
user    0m0.002s  
sys     0m0.002s
```

- **real:** total time until your program finished
- **user:** time spent executing your program
- **sys:** time the system spent on behalf of your program

Note: blunt tool



Linux **time** command

```
int main() {  
    int temp;  
    for (int i = 0; i < 1000000000; i++) {  
        temp = i;  
    }  
    return temp;  
}
```



godbolt.org Compiler Explorer:

```
main:  
    mov     eax, 999999999  
    ret
```

```
$ g++ -O3 time_example.cpp -o time_example  
$ time ./time_example
```

```
real    0m0.009s  
user    0m0.002s  
sys     0m0.002s
```

- **real:** total time until your program finished
- **user:** time spent executing your program
- **sys:** time the system spent on behalf of your program

Note: blunt tool



Compiler pitfalls

```
int main() {  
    int size = 1000;  
    int index = 0;  
    for (int i = 0; i < 1000000000; i++) {  
        index = i % size;  
        // ... vector operations  
    }  
}
```

standard string-to-int conversion

```
int main() {  
    int size = std::stoi(string_from_user);  
    int index = 0;  
    for (int i = 0; i < 1000000000; i++) {  
        index = i % size;  
        // ... vector operations  
    }  
}
```

Compiler pitfalls

```
int main() {
    int size = 1000;
    int index = 0;
    for (int i = 0; i < 1000000000; i++) {
        index = i % size;
        // ... vector operations
    }
}
```

Array size: 1000
Steps taken: 1000000000
Time per step: 1.06009 ns

standard string-to-int conversion

```
int main() {
    int size = std::stoi(string_from_user);
    int index = 0;
    for (int i = 0; i < 1000000000; i++) {
        index = i % size;
        // ... vector operations
    }
}
```

Array size: 1000
Steps taken: 1000000000
Time per step: 2.02596 ns

⇒ loop 2x slower

Compiler pitfalls

```
int main() {
    int size = 1000;
    int index = 0;
    for (int i = 0; i < 1000000000; i++) {
        index = i % size;
        // ... vector operations
    }
}
```

```
movabs rsi, 2361183241434822607
```

```
.L4:
mov rdx, rcx
shr rdx
mov rax, rdx
mul rsi
mov rax, rcx
shr rdx, 4
imul rdx, rdx, 250
sub rax, rdx
movsx rax, DWORD PTR [rbp+0+rax*4]
add ebx, DWORD PTR [rbp+0+rax*4]
```

standard string-to-int
conversion

```
int main() {
    int size = std::stoi(string_from_user);
    int index = 0;
    for (int i = 0; i < 1000000000; i++) {
        index = i % size;
        // ... vector operations
    }
}
```

```
mov rax, rcx
cqo
idiv rbp
movsx rax, DWORD PTR [rbx+rdx*4]
add esi, DWORD PTR [rbx+rax*4]
mov r12d, esi
```

⇒ loop 2x slower


Benchmarking tools

Some more sophisticated benchmarking tools:

- C/C++
 - Catch2
 - Google Benchmark
- Rust
 - bench
- Python
 - pytest-benchmark
 - timeit
- General purpose
 - hyperfine
 - **perf** (Linux)



Benchmarking - conclusions



- Compiler optimizations
- Programs running in parallel \Rightarrow variance
- CPUs boost clock frequencies
- ...

```

movabs rsi, 236118324143
.L4:
mov rdx, rcx
shr rdx
mov rax, rdx
mul rsi
mov rax, rcx
shr rdx, 4
imul rdx, rdx, 250
sub rax, rdx
movsx rax, DWORD PTR [rbp+0+rax*4]
add ebx, DWORD PTR [rbp+0+rax*4]

mov rax, rcx
cqo
idiv rbp
movsx rax, DWORD PTR [rbx+rdx*4]
add esi, DWORD PTR [rbx+rax*4]
mov r12d, esi
  
```

Healthy scepticism and reasonable expectations

Conclusions - conclusions

Tools and Techniques, Lecture 3

Performance

- Algorithmic complexity
- Memory access patterns
- Benchmarking

<https://github.com/Stoneandbeach/CSC2025>



Backup



Catch2

```
#include <catch2/catch_test_macros.hpp>
#include <catch2/benchmark/catch_benchmark.hpp>

int arithmetic_sum(int upper) {
    volatile int sum = 0;
    for (int i = 1; i <= upper; i++) {
        sum += i;
    }
    return sum;
}

TEST_CASE("Benchmarking Basics") {
    BENCHMARK("Arithmetic sum 1 to 100") {
        return arithmetic_sum(100);
    };
}
```

Catch2

```
#include <catch2/catch_test_macros.hpp>
#include <catch2/benchmark/catch_benchmark.hpp>

int arithmetic_sum(int upper) {
    volatile int sum = 0;
    for (int i = 1; i <= upper; i++) {
        sum += i;
    }
    return sum;
}

TEST_CASE("Benchmarking Basics") {
    BENCHMARK("Arithmetic sum 1 to 100") {
        return arithmetic_sum(100);
    };
}
```

including Catch2
library

function to
benchmark

defining a set of tests
or benchmarks

configuring a
benchmark

Catch2

```
#include <catch2/catch_test_macros.hpp>
#include <catch2/benchmark/catch_benchmark.hpp>

int arithmetic_sum(int upper) {
    volatile int sum = 0;
    for (int i = 1; i <= upper; i++) {
        sum += i;
    }
    return sum;
}

TEST_CASE("Benchmarking Basics") {
    BENCHMARK("Arithmetic sum 1 to 100") {
        return arithmetic_sum(100);
    };
}
```

including Catch2
library

function to
benchmark

defining a set of tests
or benchmarks

configuring a
benchmark

Catch2 output


```
-----
Benchmarking Basics
-----
```

```
/eos/user/k/kaastran/schools/CSC2025/prep/exercises/src/ex0.0_benchmarking_basics.cpp:15
.....
```

benchmark name	samples	iterations	est run time
	mean	low mean	high mean
	std dev	low std dev	high std dev
Arithmetic sum 1 to 100	100	518	3.4706 ms
	38.5609 ns	36.9933 ns	40.8075 ns
	9.47688 ns	7.24546 ns	12.492 ns

```
=====
test cases: 1 | 1 passed
assertions: - none -
```

???



Want to learn about
hardware? Pick up
electronics as a hobby!

Further reading

Latency Numbers Every Programmer Should Know (originally from Jeff Dean):

https://people.eecs.berkeley.edu/~rcs/research/interactive_latency.html

A crazy breakdown of CPU instruction latency and throughput:

https://www.agner.org/optimize/instruction_tables.pdf

Nice discussion about CPU cycle and memory access costs:

<http://ithare.com/infographics-operation-costs-in-cpu-clock-cycles/>

Modern microprocessors in 90 minutes:

<https://www.lighterra.com/papers/modernmicroprocessors/>