

Tools and Techniques, Lecture 3

Performance

- Algorithmic complexity
- Memory access patterns
- Benchmarking



LUNDS UNIVERSITET



CERN
School of Computing

CERN School of Computing 2025, Lund
Sten Åstrand, Lund University

Why?



Image credit: Giulio Eulisse, "Benchmarking and Profiling" lecture, CSC2024

How does it scale?

ALICE
Event Processing
Nodes Cluster

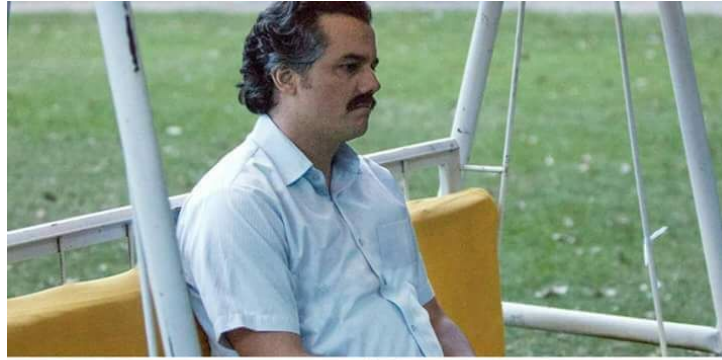
24640 CPU cores
2800 GPUs



Image credit: Giulio Eulisse, "Benchmarking and Profiling" lecture, CSC2024



How do you scale?



*Man making Python plots,
color photograph*



Image from imgflip.com

Algorithmic complexity

Theoretical metric for estimating resource consumption (in our case runtime - time complexity).

“The time complexity of an algorithm represents the number of steps it has to take to complete.”

– Complexity Theory. Brilliant.org. Retrieved 13:58, June 18, 2025, from <https://brilliant.org/wiki/complexity-theory/>

Also useful: Devopedia. 2022. "Algorithmic Complexity." Version 8, February 19. Accessed 2024-06-25. <https://devopedia.org/algorithmic-complexity>

Big O notation




⇒ algorithm is on the order of $f(n)$ or $O(f(n))$

Example: if an algorithm's runtime grows **linearly with the input size**, the **algorithm is $O(n)$**


Strictly: $O(f(n)) \Leftrightarrow$ “asymptotically bounded by $f(n)$ up to some constant”
(see Bachmann-Landau notation)

Big O Notation. Brilliant.org. Retrieved 13:54, June 18, 2025, from <https://brilliant.org/wiki/big-o-notation/>

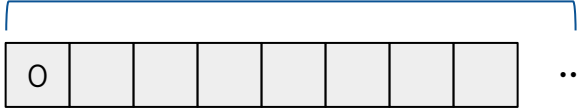
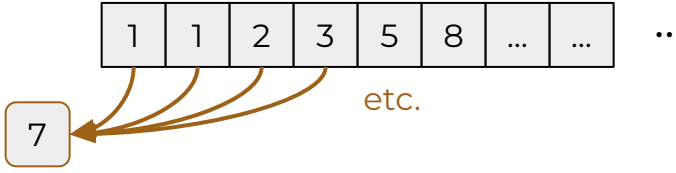
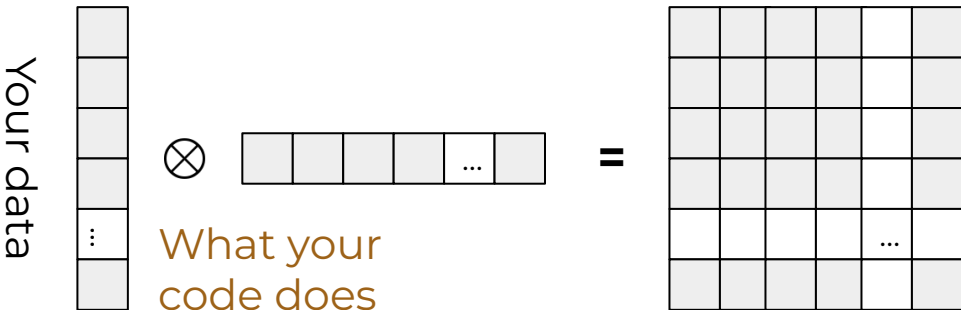
Common Big O's and examples thereof




Common Big O's and examples thereof






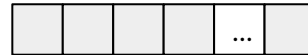
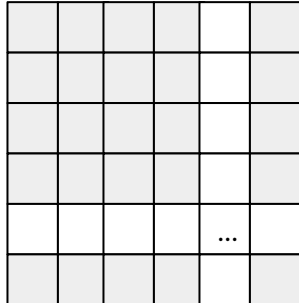
Common Big O's and examples thereof

	n elements	Complexity	Time
Your data		$O(1)$	Constant
Your data What your code does		?	?
Your data What your code does		?	?


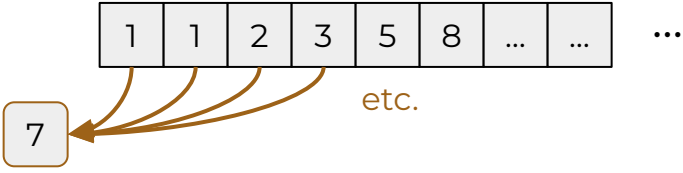
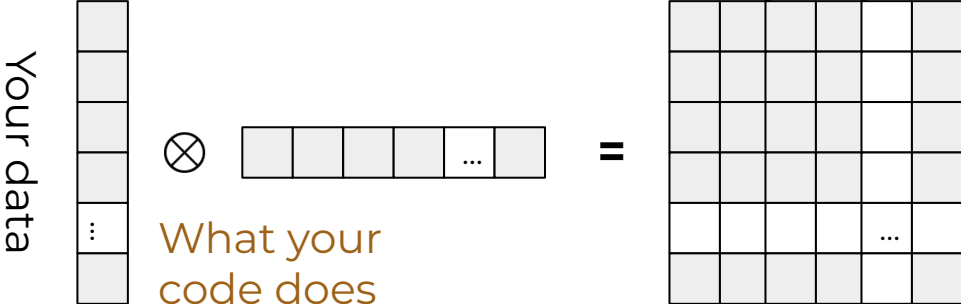
Common Big O's and examples thereof

	n elements	Complexity	Time
Your data		$O(1)$	Constant

Your data			
What your code does		$O(n)$	Linear

Your data		\otimes 	=		?	?
		What your code does				

Common Big O's and examples thereof

	n elements	Complexity	Time
Your data		$O(1)$	Constant
Your data What your code does		$O(n)$	Linear
Your data What your code does		$O(n^2)$	Quadratic

Common Big O's and examples thereof

$O(n)$ - linear time (i.e. input size doubles = runtime doubles)

- Radix sort

$O(n^2)$ - quadratic time (input size doubles = runtimes quadruples)

- “For loop in a for loop”
- Outer product of two n -length vectors $\Rightarrow n$ -by- n matrix
- Selection sort, insertion sort

Common Big O's and examples thereof

$O(n)$ - linear time (i.e. input size doubles = runtime doubles)

- Radix sort

aktshually $O(k \cdot n)$

$O(n^2)$ - quadratic time (input size doubles = runtimes quadruples)

- “For loop in a for loop”
- Outer product of two n -length vectors $\Rightarrow n$ -by- n matrix
- Selection sort, insertion sort

Common Big O's and examples thereof, cont.

$O(\log_2 n)$ - logarithmic time

- Binary search of a sorted array

$O(n \log_2 n)$ - log-linear or “linearithmic” time
(close to $O(n)$ for “reasonably sized” n)

- Many sorting algorithms, e.g. merge sort, heap sort

$O(n^3)$ - cubic time

- “For loop in a for loop in a for loop”
- “Schoolbook” matrix multiplication



Common Big O's and examples thereof, cont.

$O(\log_2 n)$ - logarithmic time

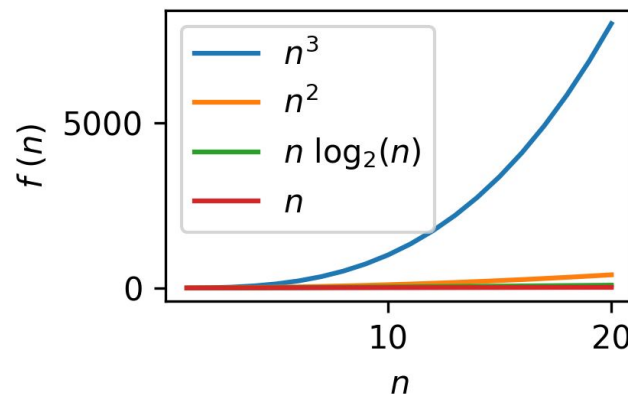
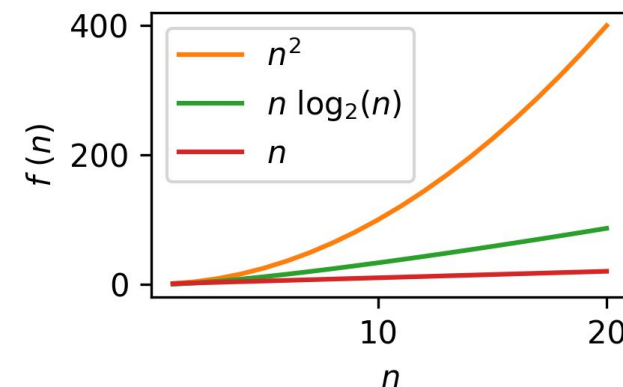
- Binary search of a sorted array

$O(n \log_2 n)$ - log-linear or “linearithmic” time
(close to $O(n)$ for “reasonably sized” n)

- Many sorting algorithms, e.g. merge sort, heap sort

$O(n^3)$ - cubic time

- “For loop in a for loop in a for loop”
- “Schoolbook” matrix multiplication



Two really Big O's (to avoid if possible)

$O(2^n)$ and **$O(n!)$** - exponential time and factorial time

- Brute-force search, combinatorics, NP-hard problems
- Finding prime numbers



Matrix multiplication

- “Schoolbook” matrix multiplication $O(n^3)$
- The Strassen algorithm (1969) $O(n^{2.805})$
- State of the art (2023) $O(n^{2.373})$

For 4 x 4 matrices:

Strassen algorithm: 49 multiplications

Google Deepmind AlphaEvolve algorithm: 48 multiplications

AlphaEvolve: A Gemini-powered coding agent for designing advanced algorithms

<https://deepmind.google/discover/blog/alphaevolve-a-gemini-powered-coding-agent-for-designing-advanced-algorithms/>





Algorithmic complexity - conclusion

```
int binary_search(const std::vector<int>& v, int target) {  
    int left = 0, right = v.size() - 1;  
  
    while (left <= right) {  
        int mid = left + (right - left) / 2;  
        if (v[mid] == target)  
            return mid;  
        else if (v[mid] < target)  
            left = mid + 1;  
        else  
            right = mid - 1;  
    }  
  
    return -1; // not found  
}
```

Common code patterns

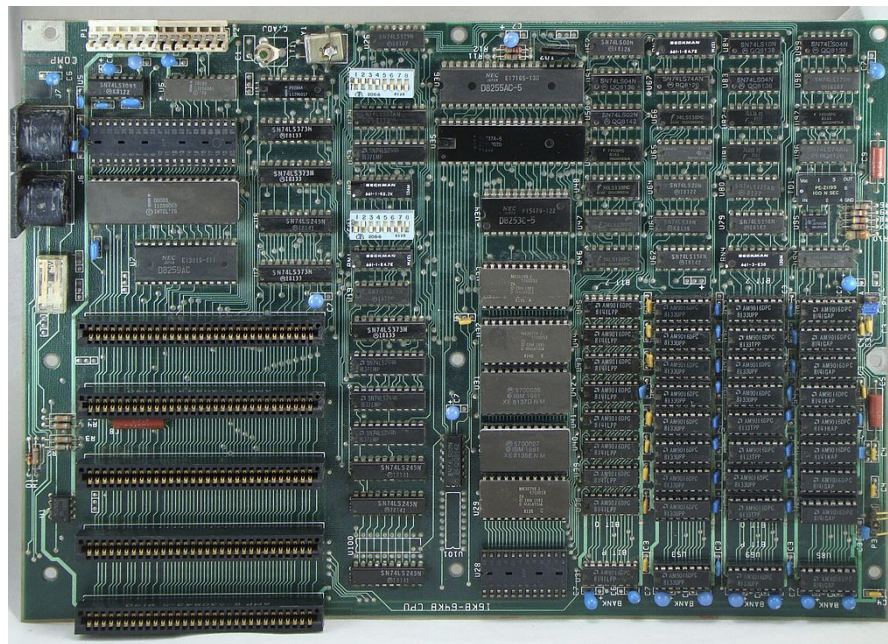
- For loop in a for loop
 - Solve differently?
 - Combine loops?
- Searching through data
 - Consider sorting first?
 - Use look-up table?
- Recomputing values
 - Compute once, keep in a table?



Pit stop



Memory access patterns



An IBM Personal Computer Model 5150 motherboard, 1981 - the first "PC". Image credit: user GermanX on Wikimedia Commons, under license [Attribution-ShareAlike 2.5 Generic](https://creativecommons.org/licenses/by-sa/2.5/)

Memory access patterns

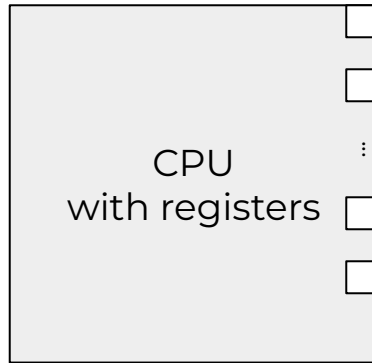
Image credit: user Gravislizard on
Wikimedia Commons, under license
[Attribution-ShareAlike 2.0 France](#)



An IBM Personal Computer Model 5150 motherboard, 1981 - the first "PC". Image credit: user GermanX on Wikimedia Commons, under license [Attribution-ShareAlike 2.5 Generic](#)

CPU-to-memory layout

CPU registers

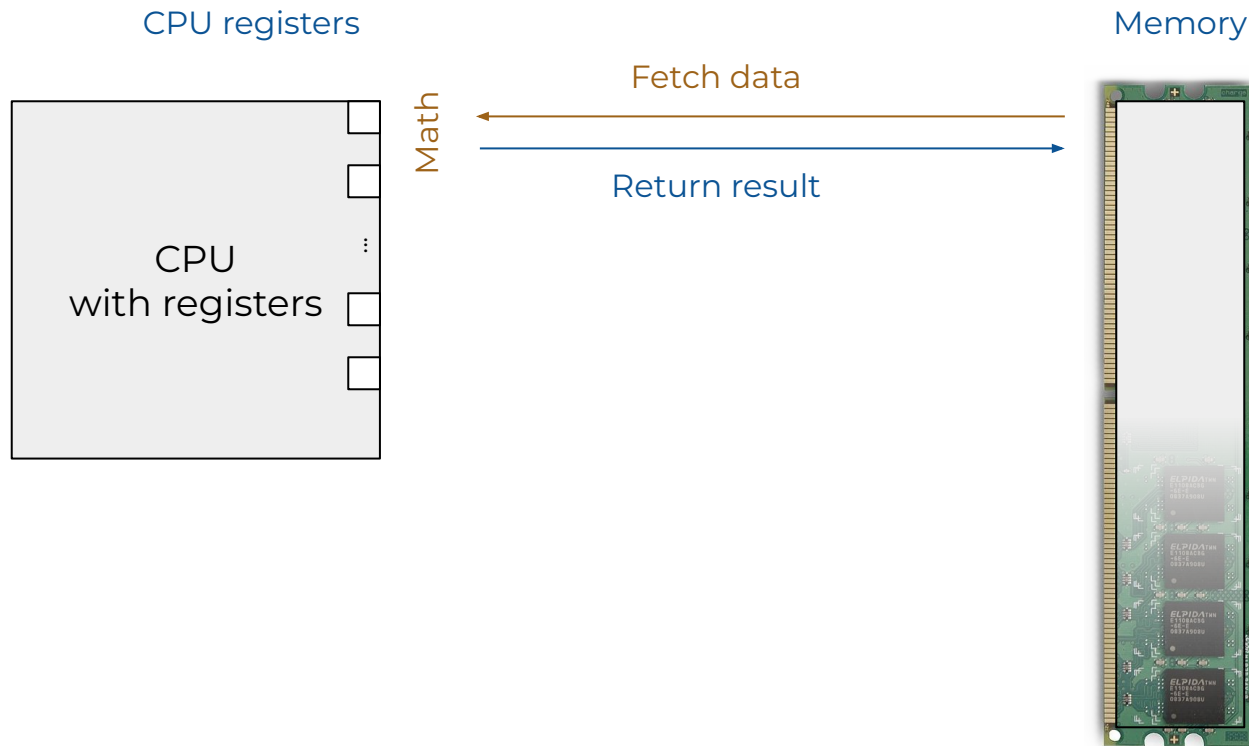


Memory



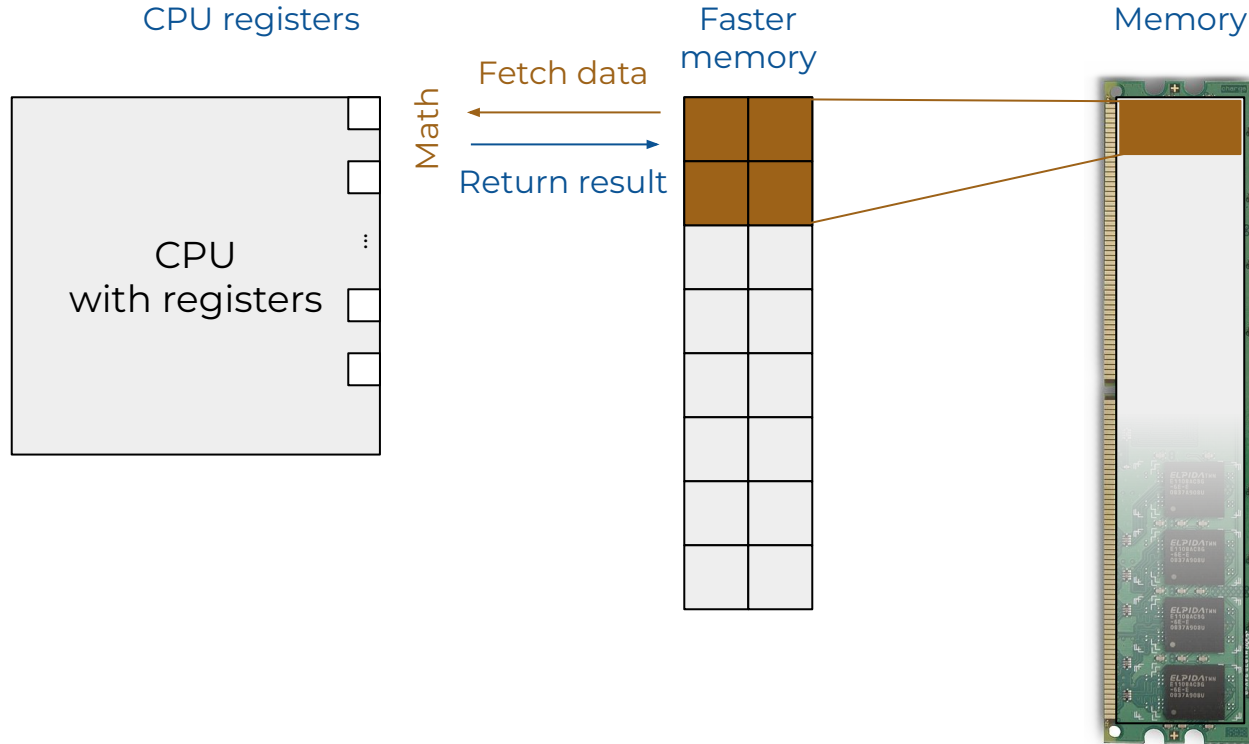
RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

CPU-to-memory layout



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

CPU-to-memory layout



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

CPU-to-memory layout

CPU registers




L1 cache

L2 cache

L3 cache

RAM

SSD




RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

CPU-to-memory layout

Typical sizes of different memory hardware



RAM
O(10 GB)



SSD
O(1 TB)

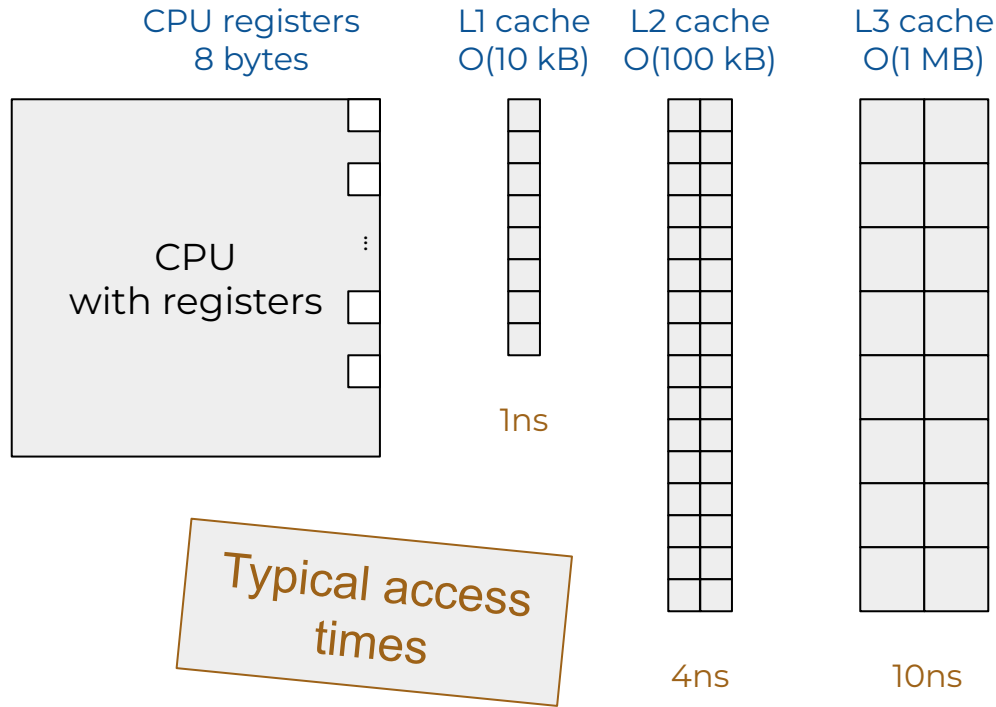


RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)


SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

CPU-to-memory layout

Typical sizes of different memory hardware



RAM
O(10 GB)



100ns

SSD
O(1 TB)




16μs

RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

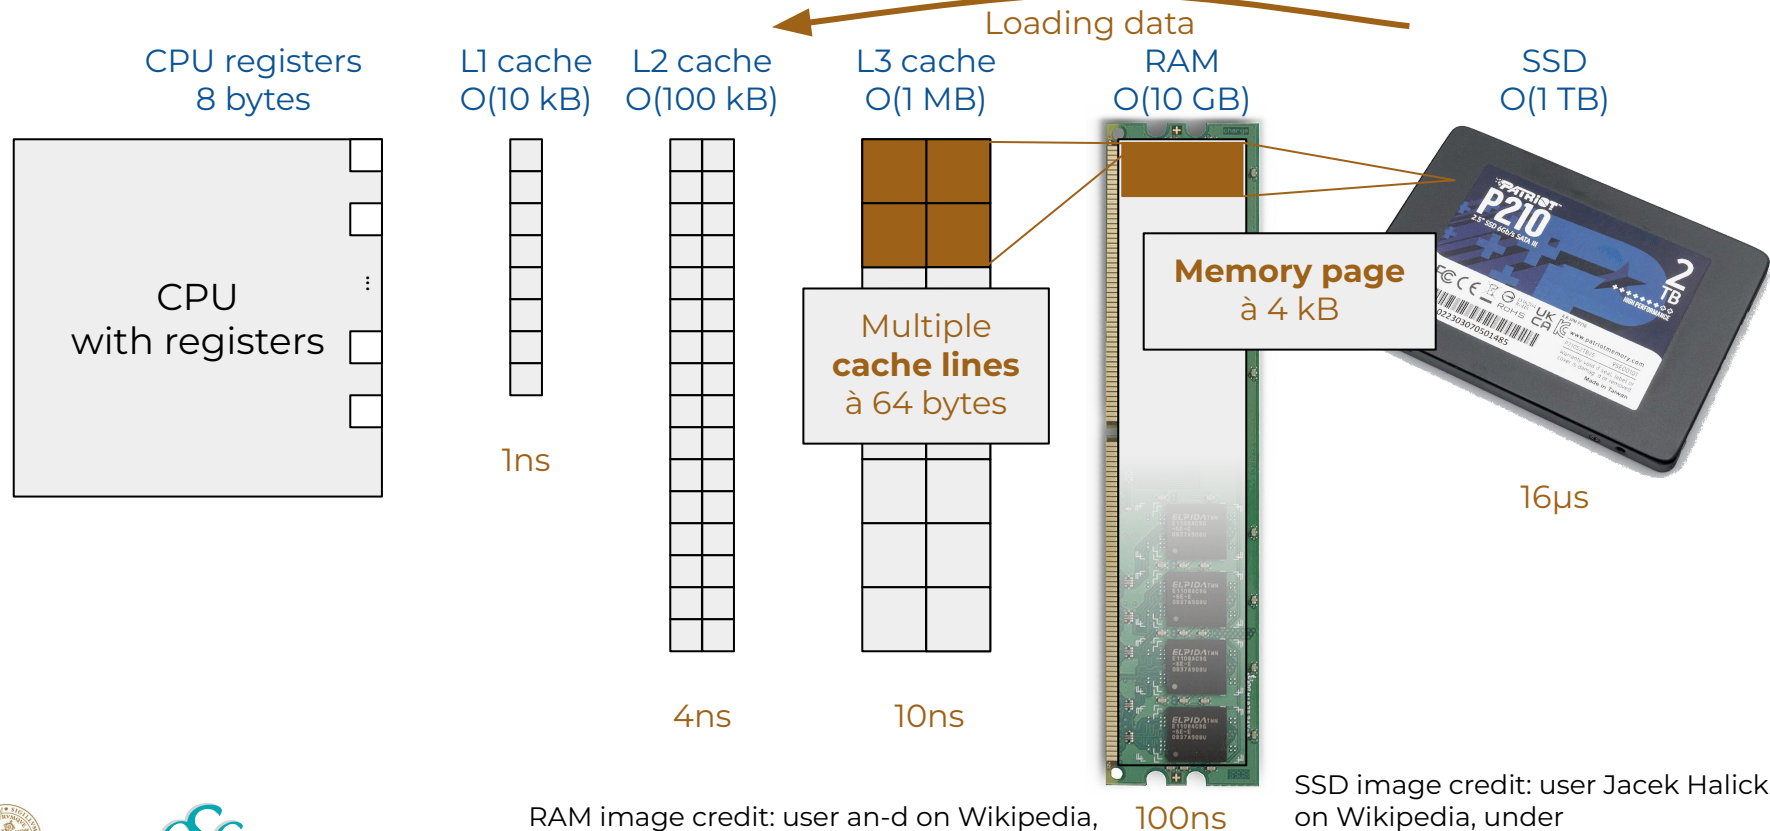
Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)


Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

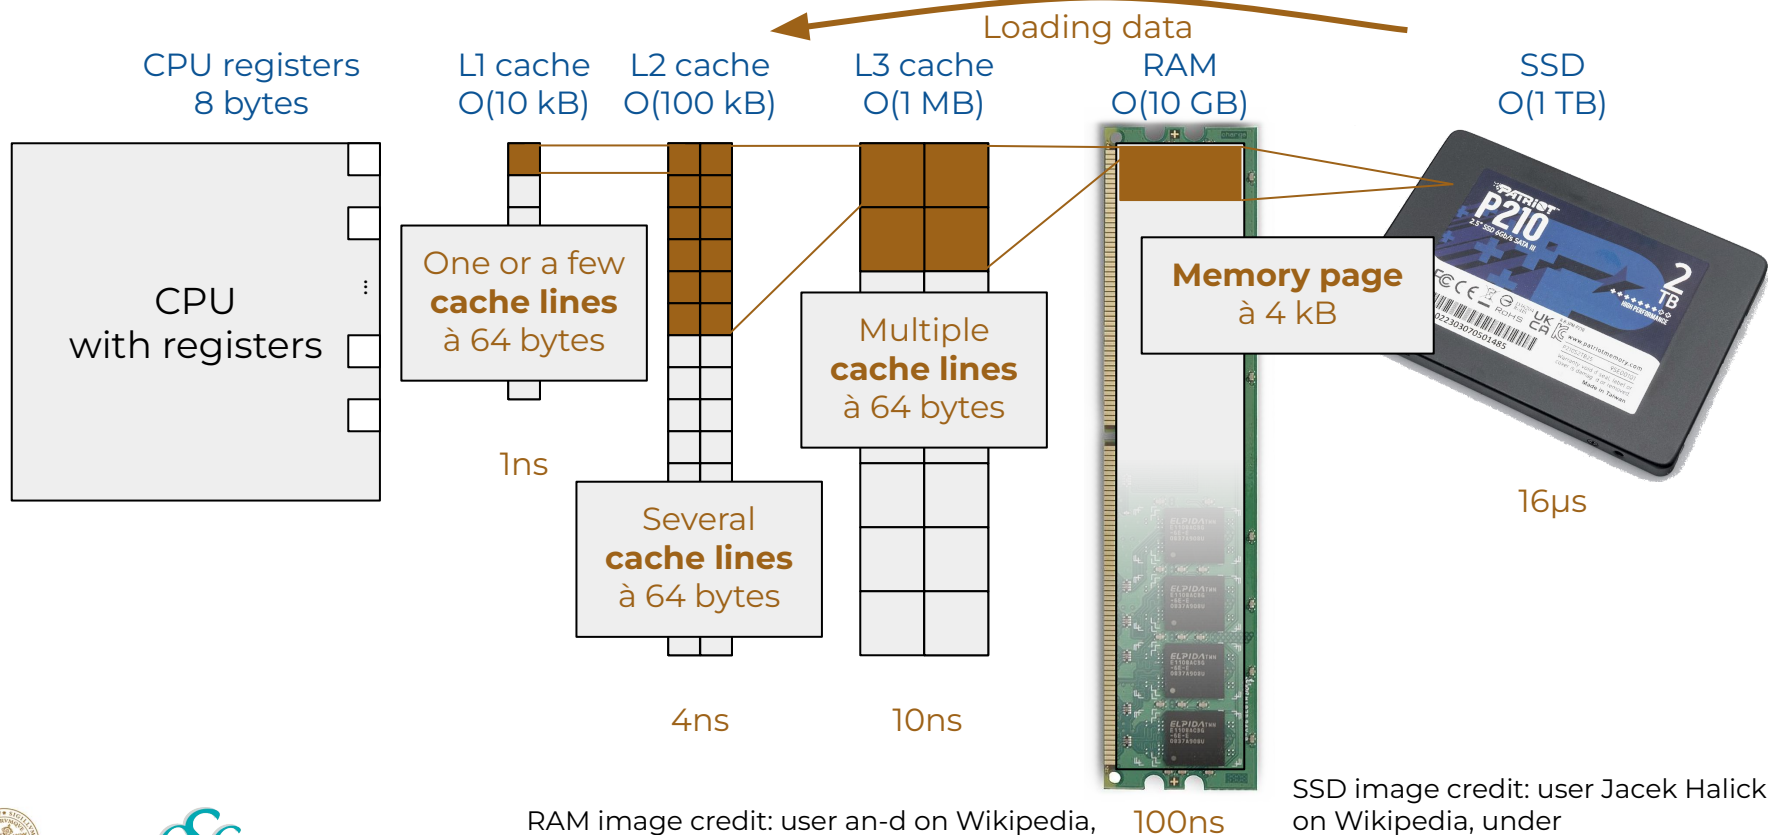
Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)


Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

Block-based memory access, the “how”



RAM image credit: user an-d on Wikipedia,
under [Attribution-ShareAlike 3.0 Unported](#)

SSD image credit: user Jacek Halicki
on Wikipedia, under
[Attribution-ShareAlike 4.0](#)

Block-based memory access, the “why”

“**Principle of locality**” or “**data locality**” - data access often happens on many elements close to each other.

If you read
this...



You will
likely then
read those

Block-based memory access, the “why”

“**Principle of locality**” or “**data locality**” - data access often happens on many elements close to each other.

If you read
this...



You will
likely then
read those


- **space locality**
- **time locality**

Pre-fetching: loading data that is likely to be needed soon

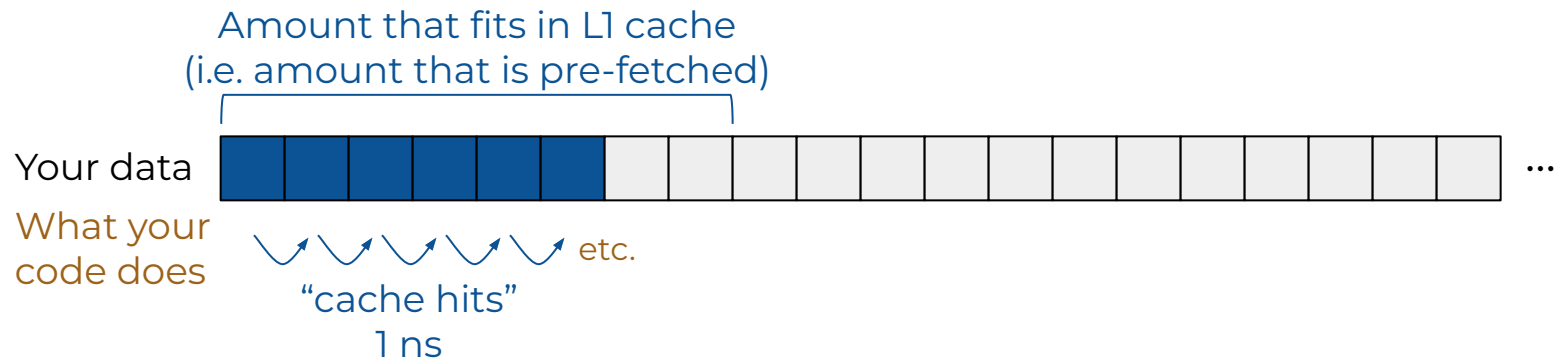
Memory access patterns

Amount that fits in L1 cache
(i.e. amount that is pre-fetched)


Your data



Memory access patterns



Memory access patterns



Memory access patterns

Amount that fits in L1 cache
(i.e. amount that is pre-fetched)

Your data



What your
code does

"cache hits"

1 ns

"cache miss"

4 ns

Your data



What your
code does

etc.

Memory access patterns

Amount that fits in L1 cache
(i.e. amount that is pre-fetched)

Your data



What your
code does




"cache hits"

1 ns


"cache miss"

4 ns

Your data




What your
code does



40




Memory access patterns - conclusions



Pit stop 2



Benchmarking



Benchmarking, pitfalls of

...but measuring is hard.

- Compiler optimizations
- Programs running in parallel \Rightarrow variance
- CPUs boost clock frequencies
- ...

Takeaway: healthy scepticism and reasonable expectations



Linux **time** command

```
int main() {  
    int temp;  
    for (int i = 0; i < 1000000000; i++) {  
        temp = i;  
    }  
    return temp;  
}
```

```
$ g++ -O3 time_example.cpp -o time_example  
$ time ./time_example
```

```
real    0m0.009s  
user    0m0.002s  
sys     0m0.002s
```

- **real:** total time until your program finished
- **user:** time spent executing your program
- **sys:** time the system spent on behalf of your program

Note: blunt tool



Linux **time** command

```
int main() {  
    int temp;  
    for (int i = 0; i < 1000000000; i++) {  
        temp = i;  
    }  
    return temp;  
}
```



godbolt.org Compiler Explorer:

```
main:  
    mov     eax, 999999999  
    ret
```

```
$ g++ -O3 time_example.cpp -o time_example  
$ time ./time_example
```

```
real    0m0.009s  
user    0m0.002s  
sys     0m0.002s
```

- **real:** total time until your program finished
- **user:** time spent executing your program
- **sys:** time the system spent on behalf of your program

Note: blunt tool



Compiler pitfalls

```
int main() {  
    int size = 1000;  
    int index = 0;  
    for (int i = 0; i < 1000000000; i++) {  
        index = i % size;  
        // ... vector operations  
    }  
}
```

standard string-to-int conversion

```
int main() {  
    int size = std::stoi(string_from_user);  
    int index = 0;  
    for (int i = 0; i < 1000000000; i++) {  
        index = i % size;  
        // ... vector operations  
    }  
}
```

Compiler pitfalls

```
int main() {  
    int size = 1000;  
    int index = 0;  
    for (int i = 0; i < 1000000000; i++) {  
        index = i % size;  
        // ... vector operations  
    }  
}
```

Array size: 1000
Steps taken: 1000000000
Time per step: 1.06009 ns

standard string-to-int
conversion

```
int main() {  
    int size = std::stoi(string_from_user);  
    int index = 0;  
    for (int i = 0; i < 1000000000; i++) {  
        index = i % size;  
        // ... vector operations  
    }  
}
```

Array size: 1000
Steps taken: 1000000000
Time per step: 2.02596 ns

⇒ loop 2x slower

Compiler pitfalls

```
int main() {
    int size = 1000;
    int index = 0;
    for (int i = 0; i < 1000000000; i++) {
        index = i % size;
        // ... vector operations
    }
}
```

```
movabs rsi, 2361183241434822607
```

```
.L4:
mov rdx, rcx
shr rdx
mov rax, rdx
mul rsi
mov rax, rcx
shr rdx, 4
imul rdx, rdx, 250
sub rax, rdx
movsx rax, DWORD PTR [rbp+0+rax*4]
add ebx, DWORD PTR [rbp+0+rax*4]
```

standard string-to-int
conversion

```
int main() {
    int size = std::stoi(string_from_user);
    int index = 0;
    for (int i = 0; i < 1000000000; i++) {
        index = i % size;
        // ... vector operations
    }
}
```

```
mov rax, rcx
cqo
idiv rbp
movsx rax, DWORD PTR [rbx+rdx*4]
add esi, DWORD PTR [rbx+rax*4]
mov r12d, esi
```

⇒ loop 2x slower


Benchmarking tools

Some more sophisticated benchmarking tools:

- C/C++
 - Catch2
 - Google Benchmark
- Rust
 - bench
- Python
 - pytest-benchmark
 - timeit
- General purpose
 - hyperfine
 - perf (Linux)



Benchmarking - conclusions



- Compiler optimizations
- Programs running in parallel \Rightarrow variance
- CPUs boost clock frequencies
- ...

```

movabs rsi, 2361183241438
.L4:
mov rdx, rcx
shr rdx
mov rax, rdx
mul rsi
mov rax, rcx
shr rdx, 4
imul rdx, rdx, 250
sub rax, rdx
movsx rax, DWORD PTR [rbp+0+rax*4]
add ebx, DWORD PTR [rbp+0+rax*4]

mov rax, rcx
cqo
idiv rbp
movsx rax, DWORD PTR [rbx+rdx*4]
add esi, DWORD PTR [rbx+rax*4]
mov r12d, esi
  
```

Healthy scepticism and reasonable expectations

Conclusions - conclusions

Tools and Techniques, Lecture 3

Performance

- Algorithmic complexity
- Memory access patterns
- Benchmarking

<https://github.com/Stoneandbeach/CSC2025>



Backup



Catch2

```
#include <catch2/catch_test_macros.hpp>
#include <catch2/benchmark/catch_benchmark.hpp>

int arithmetic_sum(int upper) {
    volatile int sum = 0;
    for (int i = 1; i <= upper; i++) {
        sum += i;
    }
    return sum;
}

TEST_CASE("Benchmarking Basics") {
    BENCHMARK("Arithmetic sum 1 to 100") {
        return arithmetic_sum(100);
    };
}
```

Catch2

```
#include <catch2/catch_test_macros.hpp>
#include <catch2/benchmark/catch_benchmark.hpp>

int arithmetic_sum(int upper) {
    volatile int sum = 0;
    for (int i = 1; i <= upper; i++) {
        sum += i;
    }
    return sum;
}

TEST_CASE("Benchmarking Basics") {
    BENCHMARK("Arithmetic sum 1 to 100") {
        return arithmetic_sum(100);
    };
}
```

including Catch2
library

function to
benchmark

defining a set of tests
or benchmarks

configuring a
benchmark

Catch2

```
#include <catch2/catch_test_macros.hpp>
#include <catch2/benchmark/catch_benchmark.hpp>
```

including Catch2
library

```
int arithmetic_sum(int upper) {
    volatile int sum = 0;
    for (int i = 1; i <= upper; i++) {
        sum += i;
    }
    return sum;
}
```

function to
benchmark

```
TEST_CASE("Benchmarking Basics") {
    BENCHMARK("Arithmetic sum 1 to 100") {
        return arithmetic_sum(100);
    };
}
```

defining a set of tests
or benchmarks

configuring a
benchmark

Catch2 output


```
-----
Benchmarking Basics
-----
```

```
/eos/user/k/kaastran/schools/CSC2025/prep/exercises/src/ex0.0_benchmarking_basics.cpp:15
.....
```

benchmark name	samples	iterations	est run time
	mean	low mean	high mean
	std dev	low std dev	high std dev
Arithmetic sum 1 to 100	100	518	3.4706 ms
	38.5609 ns	36.9933 ns	40.8075 ns
	9.47688 ns	7.24546 ns	12.492 ns

```
=====
test cases: 1 | 1 passed
assertions: - none -
```

???



Want to learn about
hardware? Pick up
electronics as a hobby!